

# Diagnostics Project:

## Supplementary Documentation and Advice

Two bits of important information:

1. The output of your *diagnose* function should be in the form of a 10x4 matrix giving the probabilities of the four illness variables being true in each of the 10 cases. The columns are ordered in the same order as the illness variables appear in the datasets.
2. The documentation states that your *learn* function should return 'a network object'. This object is never used by any code except your own: It is immediately passed to your *diagnose* function. Accordingly the form of this object is entirely up to you, and it is only called a network object because it should of course encode the information in the Bayesian network you train from the historic data.

Additionally, you might also find it speeds things up to pre-generate lots of random numbers and store them in a vector. You can then access them by an index (which you increase each time you fetch a random number. This may be considerably faster than many calls to `runif`, each generating a single random number.

This is the first year this project is being run. If you have any suggestions regarding improving documentation or the project in general, please let me know!